

ACIT 3855 – Lab 7 – Containerization

| | |
|--------------------|--|
| Instructor | Mike Mulder (mmulder10@bcit.ca) – Sets A and C Tim Guicherd (tguicherd@bcit.ca) – Set B |
| Total Marks | 10 |
| Due Dates | End of Next Class: <ul style="list-style-type: none">• March 4th for Set C• March 7th for Sets A and B |

Purpose

- Use Docker to Containerize our Microservices into self-contained images
- Use Docker Compose to deploy our Microservices

Part 1 – Add Code to Source Code Management

You can use a private (preferred) or public GitHub or GitLab repo(s) for your code.

Create one Git repository for all your services, but put each in a separate folder. The structure should look something like the following:

| <pre>/audit_log/app.py /audit_log/app_conf.yml /audit_log/log_conf.yml /audit_log/openapi.yml /deployment/docker-compose.yml /processing/app.py /processing/app_conf.yml /processing/log_conf.yml /processing/openapi.yml /receiver/app.py /receiver/app_conf.yml /receiver/log_conf.yml /receiver/openapi.yml /storage/app.py /storage/app_conf.yml /storage/log_conf.yml /storage/openapi.yml /storage/ - Also your SQLAlchemy object files and create/drop tables scripts</pre> | <table><thead><tr><th>Name</th><th>Last commit</th></tr></thead><tbody><tr><td>audit_log</td><td>Add initial code repos</td></tr><tr><td>deployment</td><td>Update compose file</td></tr><tr><td>processing</td><td>Add initial code repos</td></tr><tr><td>receiver</td><td>Add initial code repos</td></tr><tr><td>storage</td><td>Add initial code repos</td></tr></tbody></table> | Name | Last commit | audit_log | Add initial code repos | deployment | Update compose file | processing | Add initial code repos | receiver | Add initial code repos | storage | Add initial code repos |
|--|---|------|-------------|-----------|------------------------|------------|---------------------|------------|------------------------|----------|------------------------|---------|------------------------|
| Name | Last commit | | | | | | | | | | | | |
| audit_log | Add initial code repos | | | | | | | | | | | | |
| deployment | Update compose file | | | | | | | | | | | | |
| processing | Add initial code repos | | | | | | | | | | | | |
| receiver | Add initial code repos | | | | | | | | | | | | |
| storage | Add initial code repos | | | | | | | | | | | | |

Be careful not to check-in any passwords in your config files. Leave those values blank for now – we will deal with this in a later lab. *I believe one student may have had their MySQL database “hacked” in a past instance of this course because either their passwords were checked into a public GitHub repo or they used a default or simple password!*

Part 2 – Requirements and Dockerfile

For each of your Python based microservices:

Add a **requirements.txt** file that includes the name and version of each 3rd party Python package required for that service (i.e., those you had to explicitly install).

For the Receiver service you need the following in your requirements.txt file:

```
connexion==2.7.0
swagger-ui-bundle==0.0.8
```

```
requests==2.25.1
pykafka==2.4.0
```

For the Storage service:

```
connexion==2.7.0
swagger-ui-bundle==0.0.8
SQLAlchemy==1.3.22
mysql-connector-python==8.0.23
pymysql==1.0.2
pykafka==2.4.0
```

For the Processing service:

```
connexion==2.7.0
swagger-ui-bundle==0.0.8
APScheduler==3.6.3
```

For the Audit Log service:

```
connexion==2.7.0
swagger-ui-bundle==0.0.8
requests==2.25.1
pykafka==2.4.0
```

Create a **Dockerfile** that is configured as follows:

- Using a Ubuntu base image
- Installs Python and Pip on top of the base image
- Installs the dependencies from your requirements.txt file
- Runs your Connexion application (i.e., app.py) when the image is run

Here is an example:

```
FROM ubuntu:18.04
```

```
LABEL maintainer="mmulder10@bcit.ca"
```

```
RUN apt-get update -y && \
    apt-get install -y python3 python3-pip
```

```
# We copy just the requirements.txt first to leverage Docker cache
```

```
COPY ./requirements.txt /app/requirements.txt
```

```
WORKDIR /app
```

```
RUN pip3 install -r requirements.txt
```

```
COPY ./app
```

```
ENTRYPOINT [ "python3" ]
```

CMD ["app.py"]

Update the maintainer label to be your e-mail address.

Check these files into your Git repositories. Your repositories should look something like this:

```
/audit_log/app.py
/audit_log/app_conf.yml
/audit_log/log_conf.yml
/audit_log/openapi.yml
/audit_log/dockerfile
/audit_log/requirements.txt
/deployment/docker-compose.yml
/processing/app.py
/processing/app_conf.yml
/processing/log_conf.yml
/processing/openapi.yml
/processing/dockerfile
/processing/requirements.txt
/receiver/app.py
/receiver/app_conf.yml
/receiver/log_conf.yml
/receiver/openapi.yml
/receiver/dockerfile
/storage/requirements.txt
/storage/app.py
/storage/app_conf.yml
/storage/log_conf.yml
/storage/openapi.yml
/storage/dockerfile
/storage/requirements.txt
/storage/ - Also your SQLAlchemy object files and create/drop tables scripts
```

Part 3 – Building and Running Your Docker Images

Run your Kafka, Zookeeper and MySQL via Docker Compose:

cd into the deployment folder

```
docker-compose up -d
```

Build and run your Receiver service:

cd into the repo folder for the Receiver service

```
docker build -t receiver:latest .
```

Note: You should see your new image in the list of images via the command “docker images”. The REPOSITORY will be receiver and the TAG latest. It will be assigned an IMAGE ID. You can delete an image via the command “docker image rm <IMAGE ID>”.

Now run the image as a container in the background:

```
docker run -d -p 8080:8080 receiver
```

Note: The -d option runs the container in the background. The -p 8080:8080 option exposes port 8080 on the container for our service.

Repeat for the other services in each of their repo folders:

```
docker build -t storage:latest .
```

```
docker run -d -p 8090:8090 storage
```

```
docker build -t processing:latest .
```

```
docker run -d -p 8100:8100 --network="host" processing
```

Note: We need `--network="host"` to allow the Processing services to make a RESTful API call to the Storage service on localhost (i.e., <http://localhost:8090>). The other services are communicating via the Kafka message broker using the host DNS so don't need this.

```
docker build -t processing:latest .
```

```
docker run -d -p 8110:8110 audit_log
```

Note: Make sure your `app_conf.yml` files have valid values, especially for the MySQL DB username and password and Kafka. Since we are running on a single VM, we can still use localhost for the hostname.

Make sure all your services are running by running `"docker ps"`. It should look something like this:

| CONTAINER ID | IMAGE | COMMAND | CREATED | STATUS | PORTS |
|--------------|------------------------|------------------------|--------------------|-------------------|---|
| 06a6a972e9f1 | processing | "python3 app.py" | 14 seconds ago | Up 13 seconds | |
| ygin | audit_log | "python3 app.py" | About a minute ago | Up About a minute | 0.0.0.0:8110->8110/tcp |
| cfea2f8c4187 | storage | "python3 app.py" | 12 hours ago | Up 12 hours | 0.0.0.0:8090->8090/tcp |
| slett | receiver | "python3 app.py" | 12 hours ago | Up 12 hours | 0.0.0.0:8080->8080/tcp |
| 5f9bef075255 | wurstmeister/kafka | "start-kafka.sh" | 6 days ago | Up 12 hours | 0.0.0.0:9092->9092/tcp |
| ds | mysql:5.7 | "docker-entrypoint.s_" | 6 days ago | Up 13 hours | 0.0.0.0:3306->3306/tcp, 33060/tcp |
| e6bc531cf7f9 | wurstmeister/zookeeper | "/bin/sh -c 'usr/sb_" | 6 days ago | Up 12 hours | 22/tcp, 2888/tcp, 3888/tcp, 0.0.0.0:32768->2181/tcp |
| well | | | | | |
| 0c1bf115253f | | | | | |
| fka_1 | | | | | |
| d3374aca90b9 | | | | | |
| _1 | | | | | |
| 2dbc43ea6a7 | | | | | |
| ookeeper_1 | | | | | |

Make sure each service has not errors while running. Use `docker logs <container id>` to see the log output for each service.

Make sure to commit your requirements.txt and Dockerfile to your GitHub/GitLab repositories.

Part 4 – Testing your Services

On your Cloud VM, open up the following ports to allow access to your services with externally available endpoints:

- Receiver – Port 8080
- Processing – Port 8100
- Audit Log – Port 8110

Using PostMan, exercise the following endpoints on the services running on your Cloud VM. You will need to use the DNS name of your VM rather than localhost for the URLs:

- Receiver – POST an event of each type
- Processing – GET the stats
- Audit Log – GET the first event of each type

Once you are convinced they all run, stop each container before moving on to Step 5.

Part 5 – Configure and Run our Docker Images with Docker Compose

You are going to add your services to the deployment/docker_compose.yml file of your repo so you will no longer have to run each service individually.

Make sure the repos for your services and the Docker images are up-to-date on your Cloud VM (this should be the case if you finished the above steps).

Configure your microservices in docker_compose.yml as follows:

- Add a new service for each of your microservices naming it appropriately (i.e., receiver, storage, processor, audit_log)
- Set the image and ports properties for each microservice. The image is the name of the Docker image you built for a particular service (run *docker images* to see all the locally built images) and the ports should correspond to the port on which the Flask server is running for that service.
- Set the depends_on property for the microservice to kafka and/or db if it requires either or both of those services to be up and running. The depends_on property should also be set if a service depends on another service (i.e., Processing depends on Storage since it calls the GET endpoints on Storage)
- Here is a summary of what your configuration for each service should be:

| | |
|------------|--|
| receiver | ports: 8080:8080 depends on kafka (since it's a producer) |
| storage | ports: 8090:8090 depends on kafka (since it's a consumer) and db |
| processing | ports: 8100:8100 network mode is host (i.e., network_mode: "host"). It needs to access the storage service endpoint (on localhost) to get new events. depends on storage |
| audit_log | ports: 8110:8110 depends on kafka (since it's a consumer) |

Here is an example for the receiver:

```
receiver:
  image: receiver
  ports:
    - "8080:8080"
  depends_on:
    - "kafka"
```

Make sure all your services are stopped as well as zookeeper, kafka and mysql (docker-compose down).

Now bring all the services up as follows with Docker Compose:

```
docker-compose up -d (note the -d runs the services in the background)
```

Test your services running on your VM with Postman to make sure they still all work. Troubleshoot as necessary.

Also test with jMeter. Make sure to update the hostname of the receiver service in your test case. You will be asked to demonstrate that your services are correctly running on your Cloud VM by running your jMeter load test against them.

Grading and Submission

Submit the following to the Lab 7 Dropbox on D2L:

- Your docker-compose.yml file

Demonstrate the following to your instructor:

| | |
|--|---------|
| Demonstrate that you can run each service individually with a docker run command and view the console or log output. | 4 marks |
| Demonstrate that you can run all your services with Docker Compose. Exercise the following: <ul style="list-style-type: none">• Your jMeter test against those services• GET of your /stats endpoint• GET of one or both of your Audit log endpoints | 6 marks |