



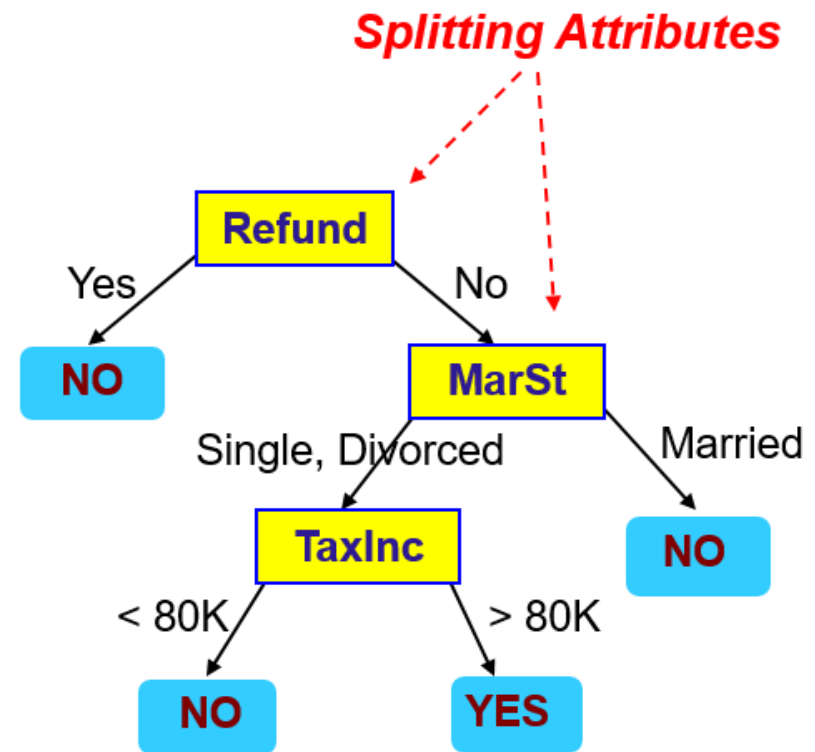
DECISION TREES



Decision Tree

<i>categorical</i>	<i>categorical</i>	<i>continuous</i>	<i>class</i>	
<i>Tid</i>	<i>Refund</i>	<i>Marital Status</i>	<i>Taxable Income</i>	<i>Cheat</i>
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

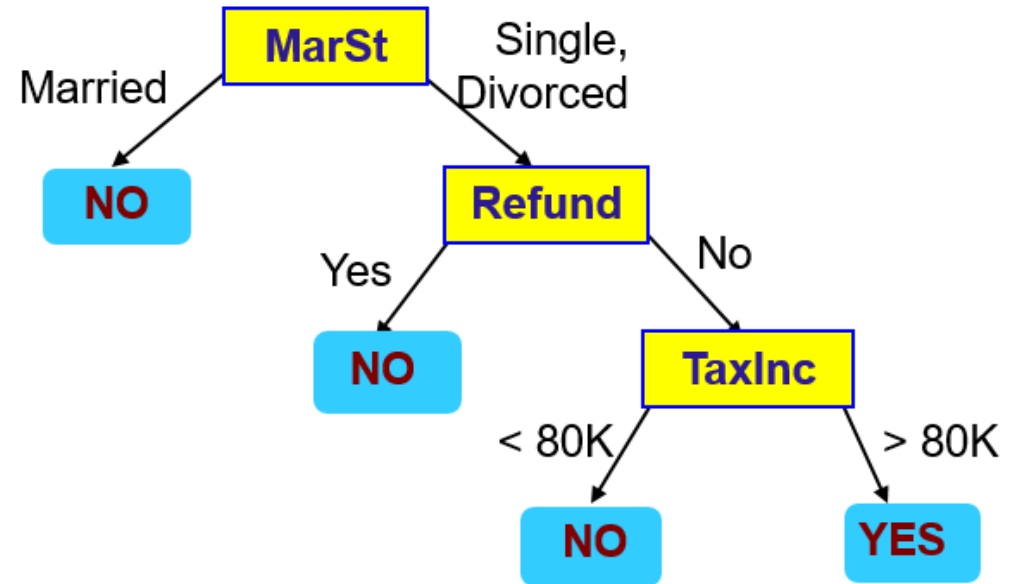
Training Data



Model: Decision Tree

Another Tree

<i>Tid</i>	<i>Refund</i>	<i>Marital Status</i>	<i>Taxable Income</i>	<i>Cheat</i>
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes



There could be more than one tree that fits the same data!

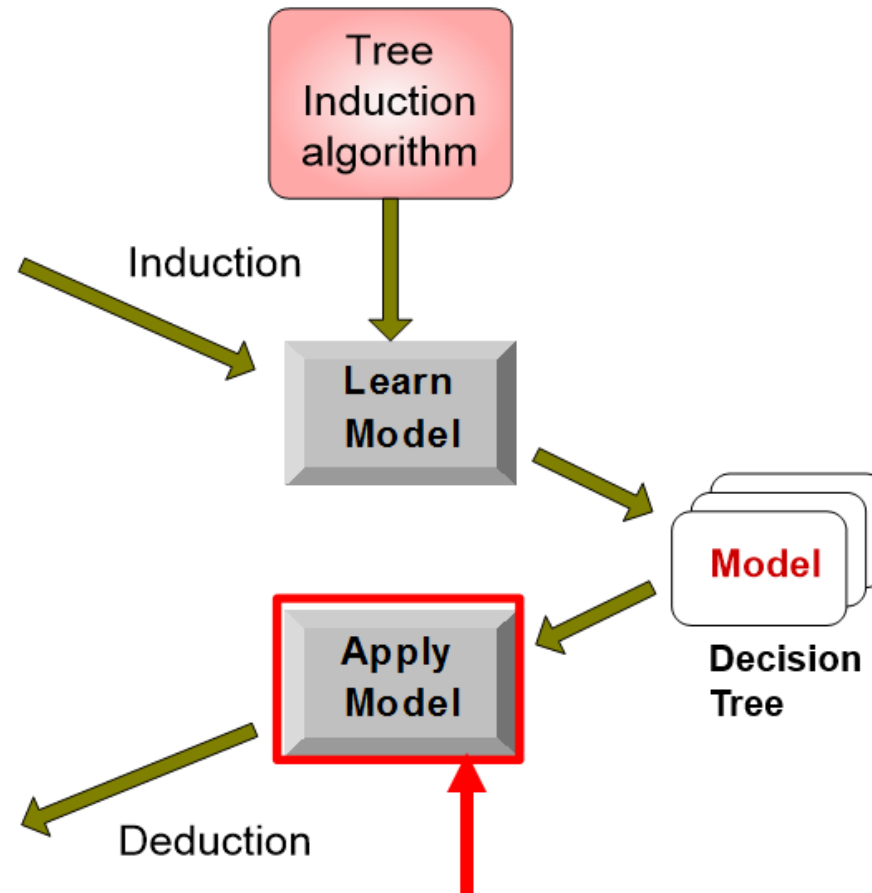
Use of Decision Tree

Tid	Attrib1	Attrib2	Attrib3	Class
1	Yes	Large	125K	No
2	No	Medium	100K	No
3	No	Small	70K	No
4	Yes	Medium	120K	No
5	No	Large	95K	Yes
6	No	Medium	60K	No
7	Yes	Large	220K	No
8	No	Small	85K	Yes
9	No	Medium	75K	No
10	No	Small	90K	Yes

Training Set

Tid	Attrib1	Attrib2	Attrib3	Class
11	No	Small	55K	?
12	Yes	Medium	80K	?
13	Yes	Large	110K	?
14	No	Small	95K	?
15	No	Large	67K	?

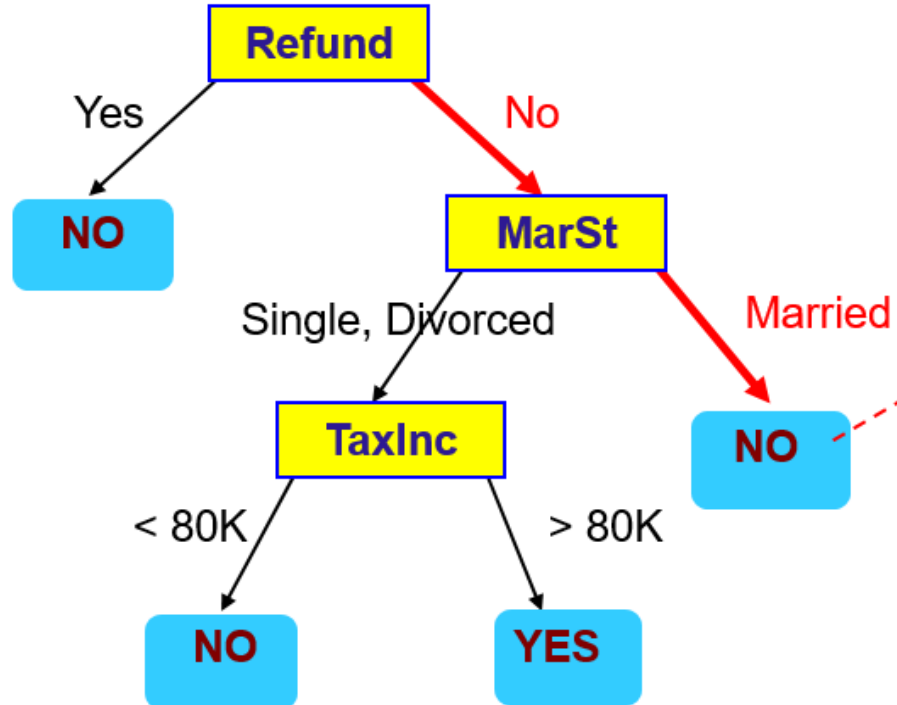
Test Set



Predict

Test Data

Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?






Assign Cheat to "No"

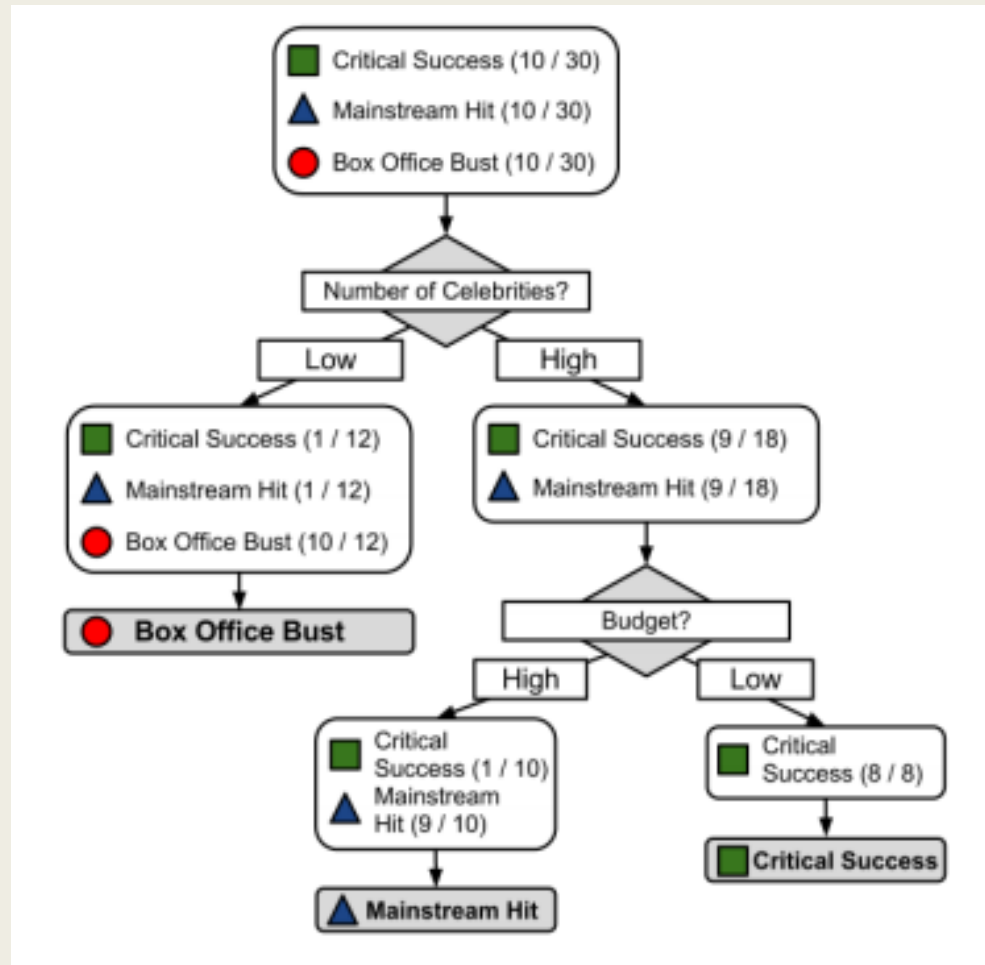
Divide and Conquer

- Decision trees are built using a heuristic called recursive partitioning, also known as divide and conquer
- Start with the root node, which is the entire dataset, choose a feature that is the most predictive of the target class
- Continue choosing the next best candidate until a stopping criteria is reached:
 - *All (or nearly all) of the nodes have the same class*
 - *No remaining features to distinguish among examples*
 - *The tree has grown to a predefined size limit*

Example: Movie Releases



Decision Tree



C5.0 Algorithm

- Uses **entropy** for measuring purity
- Entropy – indicates how mixed the class values are (0 – completely homogenous, $\log_2 c$ – maximum amount of disorder with c classes)

$$Entropy(S) = \sum_{i=1}^c -p_i \log_2(p_i)$$

- where p_i is the proportion of values falling into class i
- E.g. if partitioning a set results in 60% in one class, and 40% in another class, then entropy = $-0.6 * \log_2(0.6) - 0.4 * \log_2(0.4) = 0.97$

Definition

Entropy

lack of order or predictability; gradual decline into disorder

Information Gain

- Select two features to split (e.g. refund, marital status)
- If a split results in a higher information gain due to feature X than feature Y, then use X

$$InfoGain(F) = Entropy(S_1) - Entropy(S_2)$$

- where S_1 is the segment before the split, and the partitions resulting from the split S_2




Multiple Partitions

- Given n partitions resulting from a split S , the entropy of the split is:


$$Entropy(S) = \sum_{i=1}^n w_i Entropy(P_i)$$

- where w_i is the proportion of examples falling in that partition.

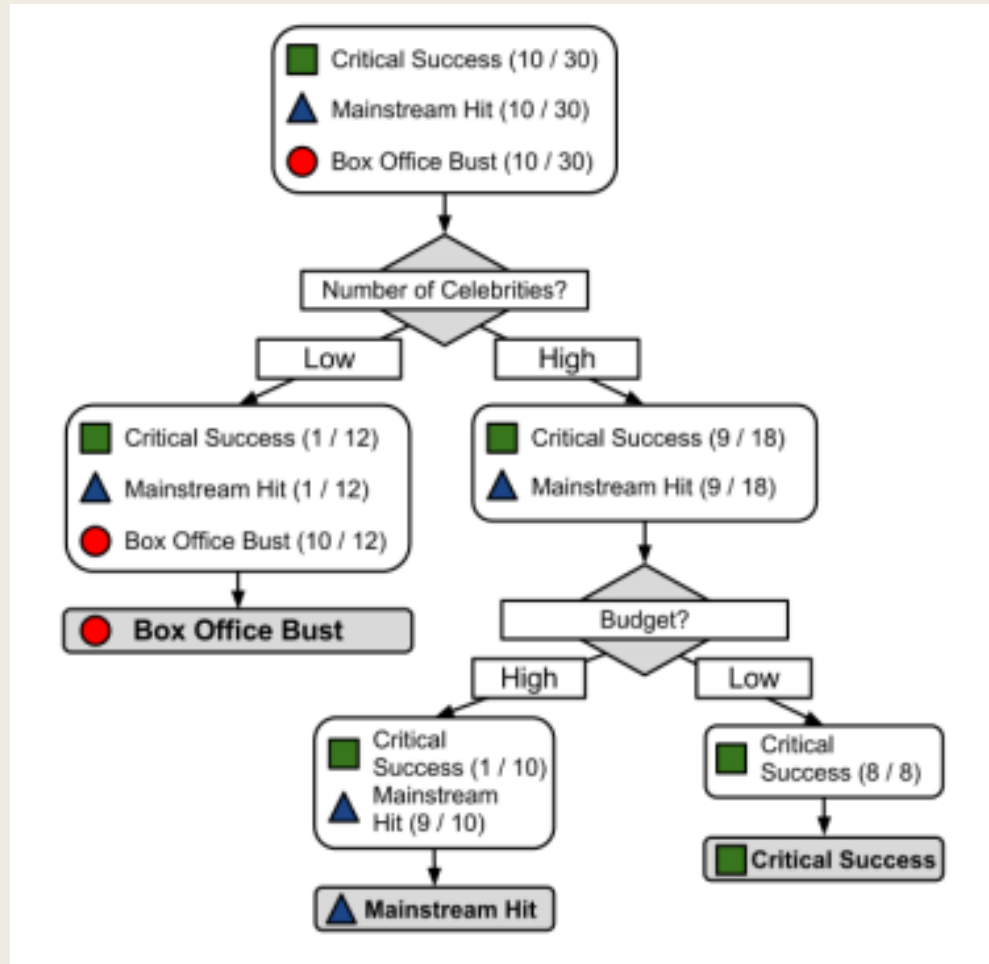
Example: Movie Releases



Decision Tree



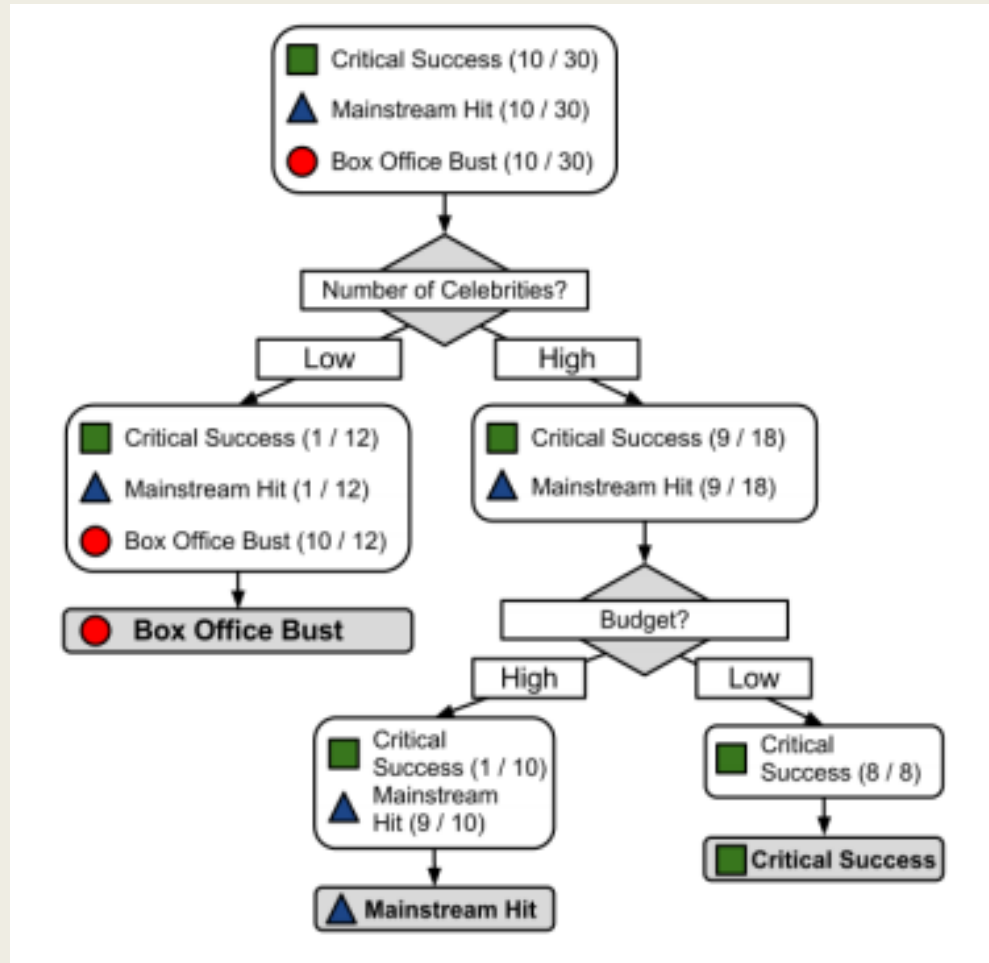
Decision Tree



$$-\frac{10}{30} \log_2\left(\frac{10}{30}\right) - \frac{10}{30} \log_2\left(\frac{10}{30}\right) - \frac{10}{30} \log_2\left(\frac{10}{30}\right)$$

1.5849

Decision Tree



$$-\frac{10}{30} \log_2\left(\frac{10}{30}\right) - \frac{10}{30} \log_2\left(\frac{10}{30}\right) - \frac{10}{30} \log_2\left(\frac{10}{30}\right)$$

1.5849




$$\frac{12}{30} * \left[-\frac{1}{12} \log_2\left(\frac{1}{12}\right) - \frac{1}{12} \log_2\left(\frac{1}{12}\right) - \frac{10}{12} \log_2\left(\frac{10}{12}\right) \right]$$

+


$$\frac{18}{30} * \left[-\frac{9}{18} \log_2\left(\frac{9}{18}\right) - \frac{9}{18} \log_2\left(\frac{9}{18}\right) \right]$$

0.5667


Alternative Split



What is the Entropy?



What is the Entropy?





$$\frac{14}{30} * \left[-\frac{1}{14} \log_2 \left(\frac{1}{14} \right) - \frac{3}{14} \log_2 \left(\frac{3}{14} \right) - \frac{10}{14} \log_2 \left(\frac{10}{14} \right) \right]$$

+



$$\frac{16}{30} * \left[-\frac{7}{16} \log_2 \left(\frac{7}{16} \right) - \frac{9}{16} \log_2 \left(\frac{9}{16} \right) \right]$$

0.757



Split Option 1



Split Option 2





Split Option 1



0.5667

Split Option 2



0.757

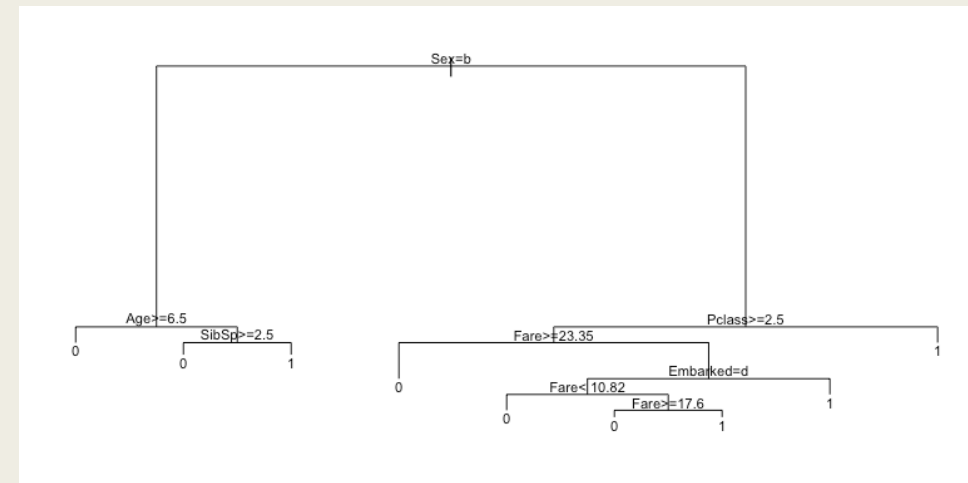
C5.0 in R

- `install.packages('C50')`
- `library(C50)`
- `model <- C5.0(train, class, trials)`
 - *train* is the training data frame without the classification
 - *class* is a factor vector with the classification for each row in *train*
 - *trials* controls the number of decision trees to be created
- `p <- predict(model, test, type="class")`
 - *model* is the model created by the C5.0 function
 - *test* is the test data data frame with the same feature as the training data frame
 - R
 - returns a vector of predicted class

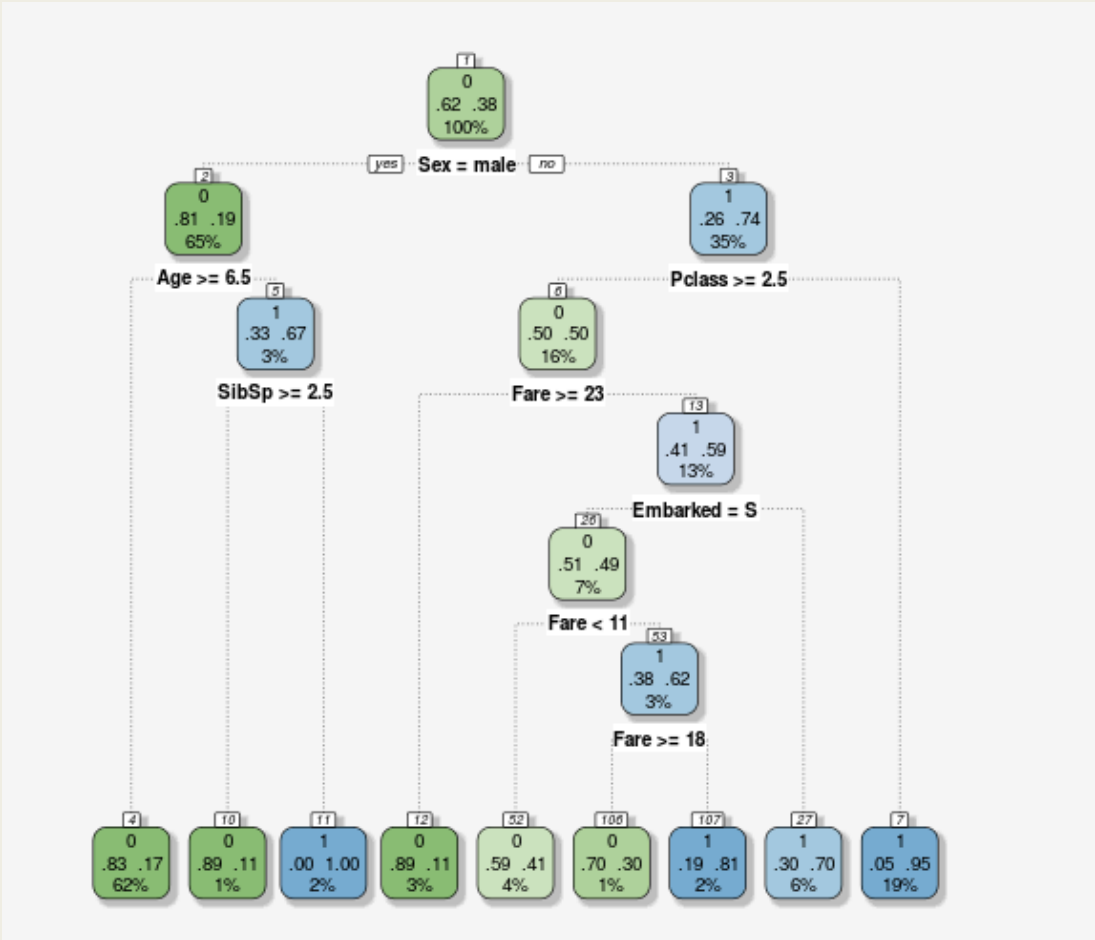
rpart

- `install.packages('rpart')`
- `library(rpart)`
- `fit <- rpart(default ~ ., data=credit, method="class")`
- `plot(fit)`
- `text(fit, cex=0.7)`

- Clear plot:
- `dev.off()`



- `install.packages('rattle')`
- `install.packages('rpart.plot')`
- `install.packages('RColorBrewer')`
- `library(rattle)`
- `library(rpart.plot)`
- `library(RColorBrewer)`
- `fancyRpartPlot(fit)`
- `pdf("tree.pdf")`
- `fancyRpartPlot(fit)`
- `dev.off()`



Predict

- `Prediction <- predict(fit, credit_test, type = "class")`
- `submit <- data.frame(Amount = credit_test$amount, Default = Prediction)`